




13. Red-Black Tree

Yu-Shuen Wang, CS, NCTU

Red-black trees

One of many search-tree schemes that are “**balanced**” in order to guarantee that basic dynamic-set operations take $O(\lg n)$ time in the worse case.



13.1 Properties of red-black tree

A **red-black tree** is a binary search tree with one extra bit of storage per node: its **color**, which can either RED or BLACK. By constraining the way nodes can be colored on any path from the root to a leaf, red-black trees ensure that no such path is more than twice as long as any other, so that three is approximately **balanced**.

13.1 Properties of red-black tree

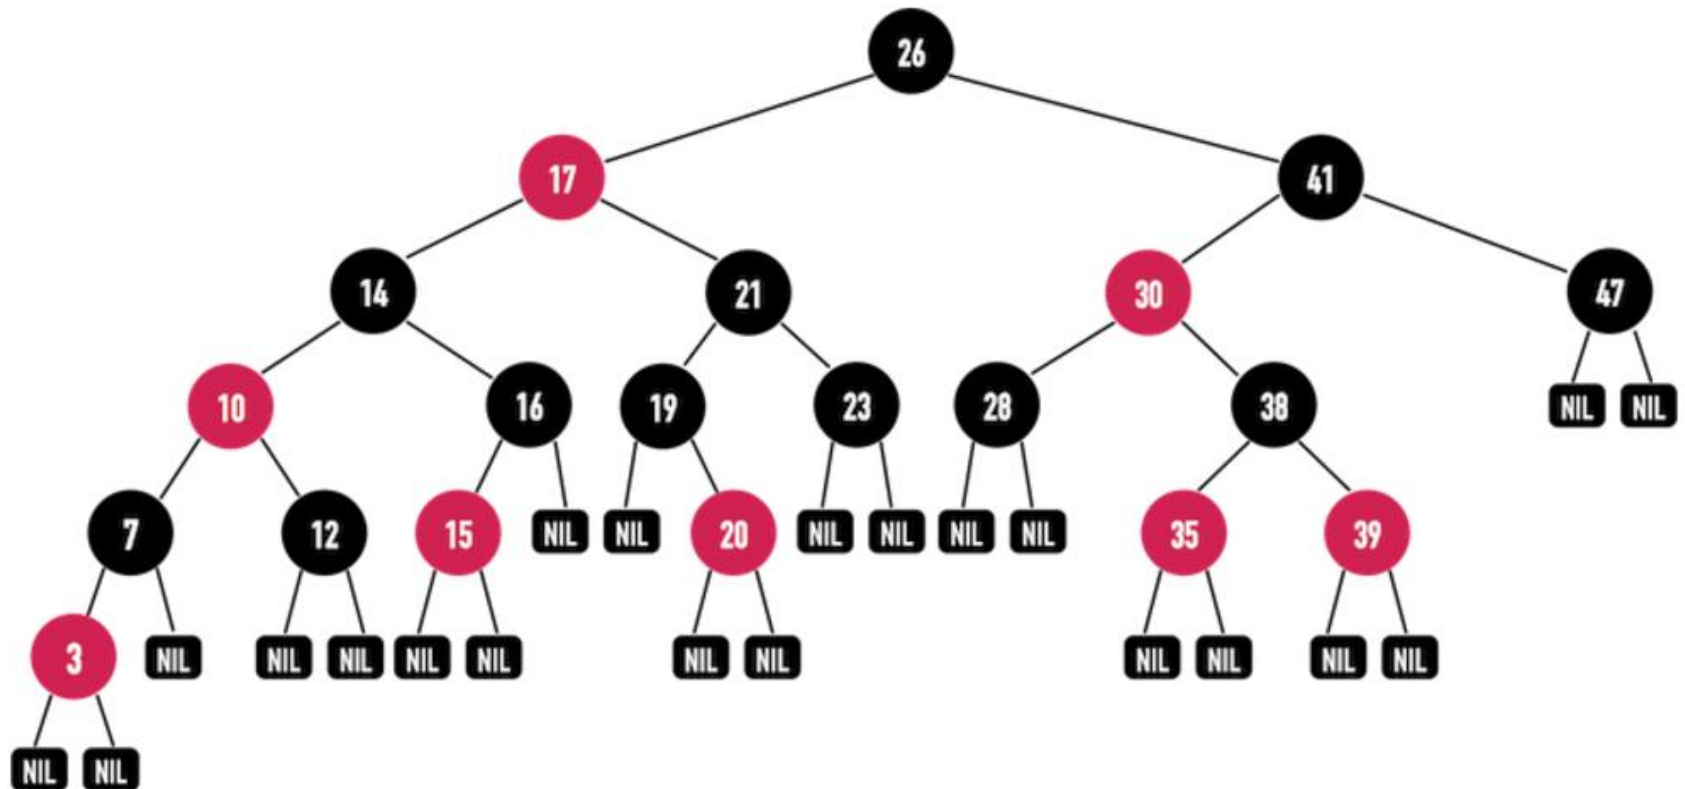
Each node of the tree now contains the fields *color*, *key*, *left*, *right*, and *p*. If a child or the parent of a node does not exist, the corresponding pointer field of the node contains the value NIL. We shall regard these NIL's as being pointers to external nodes (leaves) of the binary search tree and the normal, key-bearing nodes as being internal nodes of the tree.

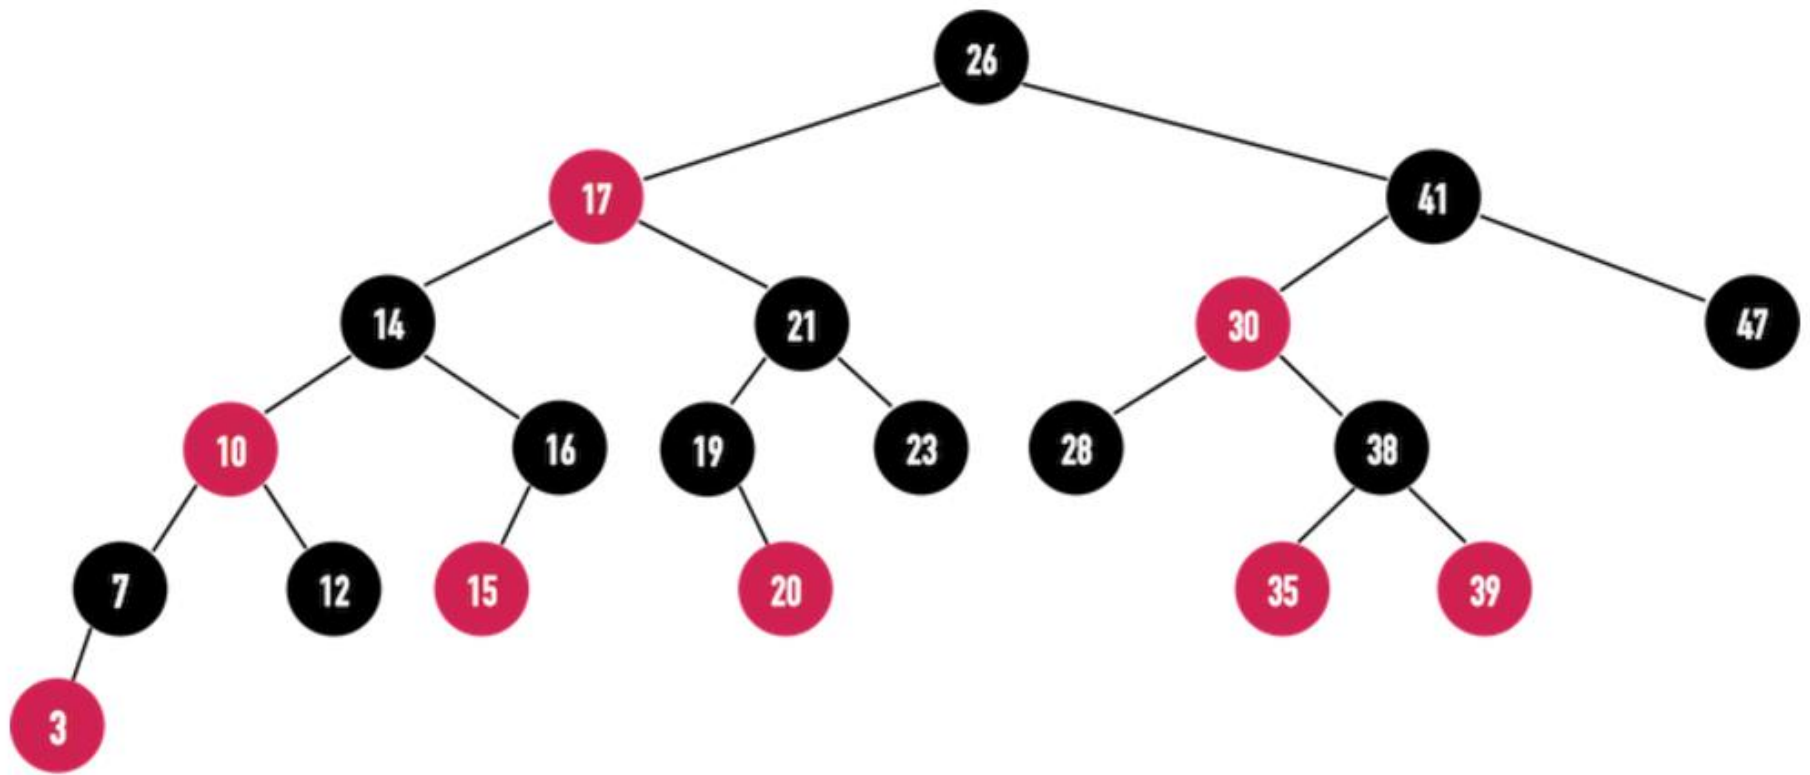
13.1 Properties of red-black tree

A binary search tree is a red-black tree if it satisfies the following **red-black properties**:

1. Every node is either red or black.
2. The root is black.
3. Every leaf (NIL) is black
4. If a node is red, then both its children are black.
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes.

A red-black tree






Lemma 13.1

- A red-black tree with n internal nodes has height at most $2\lg(n+1)$.

Proof.



- We start by showing that the subtree rooted at any node x contains at least $2^{bh(x)} - 1$ internal nodes. We prove this claim by induction on the height of x . If the height of x is 0, then x must be a leaf ($NIL[T]$), and the sub tree rooted at x indeed contains at least $2^{bh(x)} - 1 = 2^0 - 1 = 0$ internal nodes.


For the inductive step, consider a node x that has positive height and is an internal node with two children. Each child has a black-height of either $bh(x)$ or $bh(x) - 1$, depending on whether its color is red or black, respectively. Since the height of a child of x is less than the height of x itself, we can apply the inductive hypothesis to conclude that each child has at least $2^{bh(x)-1} - 1$ internal nodes. Thus, the subtree rooted at x contains at least $(2^{bh(x)-1} - 1) + (2^{bh(x)-1} - 1) + 1 = 2^{bh(x)} - 1$ internal nodes, which proves the claim.

To complete the proof of the lemma, let h be the height of the tree. According to **property 4**, at least half the nodes on any simple path from the root to a leaf, not including the root, must be black. Consequently, the black-height of the root must be at least $h/2$; thus,



$$n \geq 2^{h/2} - 1.$$

Moving the **1** to the left-hand side and taking logarithms on both sides yields $\lg(n + 1) \geq h/2$, or $2\lg(n+1) \geq h$.


13.2 Rotations – why?




13.2 Rotations



An example of LEFT-ROTATE(T, x)






LEFT-ROTATE(T, x)

```

11  1  $y \leftarrow \text{right}[x]$ 
&  2  $\text{right}[x] \leftarrow \text{left}[y]$ 
14  3  $p[\text{left}[y]] \leftarrow x$ 
    4  $p[y] \leftarrow p[x]$ 
    5 If  $p[x] = \text{nil}[T]$ 
    7   then  $\text{root}[T] \leftarrow y$ 
&  7   else if  $x = \text{left}[p[x]]$ 
18  8       then  $\text{left}[p[x]] \leftarrow y$ 
    9       else  $\text{right}[p[x]] \leftarrow y$ 
11 10  $\text{left}[y] \leftarrow x$ 
& 18 11  $p[x] \leftarrow y$ 

```



The code for RIGHT-ROTATE is **symmetric**. Both LEFT-ROTATE and RIGHT-ROTATE run in **$O(1)$** time.

13.3 Insertion

RB-INSERT(T, z)

1 $y \leftarrow nil[T]$

2 $x \leftarrow root[T]$

3 **while** $x \neq nil[T]$

► unsuccessful search

4 **do** $y \leftarrow x$

5 **if** $key[z] < key[x]$

6 **then** $x \leftarrow left[x]$

7 **else** $x \leftarrow right[x]$


8 $p[z] \leftarrow y$

13.3 Insertion

```
9  if  $y = nil[T]$                                 ► tree T was empty
10   then  $root[T] \leftarrow z$ 
11   else if  $key[z] < key[y]$                       ► link to child
12       then  $left[y] \leftarrow z$ 
13       else  $right[y] \leftarrow z$ 
14    $left[z] \leftarrow nil[T]$ 
15    $right[z] \leftarrow nil[T]$ 
16    $color[z] \leftarrow RED$ 
17    $RB-INSERT-FIXUP(T, z)$ 
```

new in
RB-tree

RB-INSERT-FIXUP



Case 1: uncle is red


Case 2: uncle is black, inserted node is right(X)

Case 3: uncle is black, inserted node is left(X)

White means the node could be red or black.


PS: uncle can be nil (black)

RB-INSERT-FIXUP – case 1




- parent turn black
- uncle turn black
- grandparent turn red
- current -> grandparent(current)
- check z (current) and w, and fix up iteratively.

RB-INSERT-FIXUP – case 2



If A is the newly inserted node, then uncle must be NIL. If not, uncle could be an ordinary node.


RB-INSERT-FIXUP – case 2



Make case 2 become case 3!


- current -> parent(current)
- left rotate at current (X)

RB-INSERT-FIXUP – case 3



Similar to case 2, if A is the newly inserted node, then uncle must be NIL. If not, uncle could be an ordinary node.


RB-INSERT-FIXUP – case 3




Why does it work?


- parent (X) turn black
- grandparent (Z) turn red
- right rotate at grandparent (Z)

In the beginning, let $bh(A)$ be M , then $bh(Z)$ is $M+1$. After the rotation, both $bh(A)$ and $bh(Z)$ are M .





case 1






case 2





case 3





RB-INSERT-FIXUP(T, z)

```
1  while  $color[p[z]] = \text{RED}$ 
2      do if  $p[z] = \text{left}[p[p[z]]]$ 
3          then  $y \leftarrow \text{right}[p[p[z]]]$ 
4              if  $color[y] = \text{RED}$ 
5                  then  $color[p[z]] \leftarrow \text{BLACK}$  Case 1
6                       $color[y] \leftarrow \text{BLACK}$  Case 1
7                       $color[p[p[z]]] \leftarrow \text{RED}$  Case 1
8                       $z \leftarrow p[p[z]]$  Case 1
9              else if  $z = \text{right}[p[z]]$ 
10                  then  $z \leftarrow p[p[z]]$  Case 2
11                       $\text{LEFT-ROTATE}(T, z)$  Case 2
12                       $color[p[z]] \leftarrow \text{BLACK}$  Case 3
13                       $color[p[p[z]]] \leftarrow \text{RED}$  Case 3
14                       $\text{RIGHT-ROTATE}(T, p[p[z]])$  Case 3
15                  else (same as then clause with “right” and “left” exchanged)
16   $color[\text{root}[T]] \leftarrow \text{BLACK}$ 
```



Analysis

- RB-INSERT take a total of $O(\lg n)$ time.
- It never performs more than two rotations, since the **while** loop terminates if case 2 or case 3 executed.

13.4 Deletion



Rule 2: the root is black.



Rule 4: parent & child shouldn't be red simultaneously





Rule 5: $bh(x)$ of all paths from x to leaves are the same.

RB-DELETE(T, z)

```
1  if  $left[z] = nil[T]$  or  $right[z] = nil[T]$ 
2      then  $y \leftarrow z$ 
3      else  $y \leftarrow \text{TREE-SUCCESSOR}(z)$ 
4  if  $left[y] \neq nil[T]$ 
5      then  $x \leftarrow left[y]$ 
6      else  $x \leftarrow right[y]$ 
7   $p[x] \leftarrow p[y]$ 
8  if  $p[y] = nil[T]$ 
9      then  $root[T] \leftarrow x$ 
10 else if  $y = left[p[y]]$ 
11     then  $left[p[y]] \leftarrow x$ 
12     else  $right[p[y]] \leftarrow x$ 
13 if  $y \neq z$ 
14 then  $key[z] \leftarrow key[y]$ 
15     copy  $y$ 's satellite data into  $z$ 
16 if  $color[y] = \text{BLACK}$ 
17     then RB-DELETE-FIXUP( $T, x$ )
18 return  $y$ 
```

- ▶ one or no child
- ▶ two children
- ▶ set x to be y 's child
- ▶ connect the child to its parent
- ▶ y is root
- ▶ y will be deleted, x becomes root
- ▶ connect parent to child

RB-DELETE-FIXUP




Delete A

Case 1. sibling is red

Case 2. sibling is black, and its children are black


Case 3. sibling is black, and its left child is red, right child is black

Case 4. sibling is black, and its right child is red





1. sibling turns black
2. current->parent turns red
3. left rotation at current->parent
4. update sibling's position

Since current node (A) is still black, the problem still exists, and will become case 2, 3, or 4.





1. sibling turns red
2. $\text{current} = \text{current} \rightarrow \text{parent}$

current (c) could be red or black after the operation






continue from here !



1. current (c) turns black

If current (c) is red, we need to

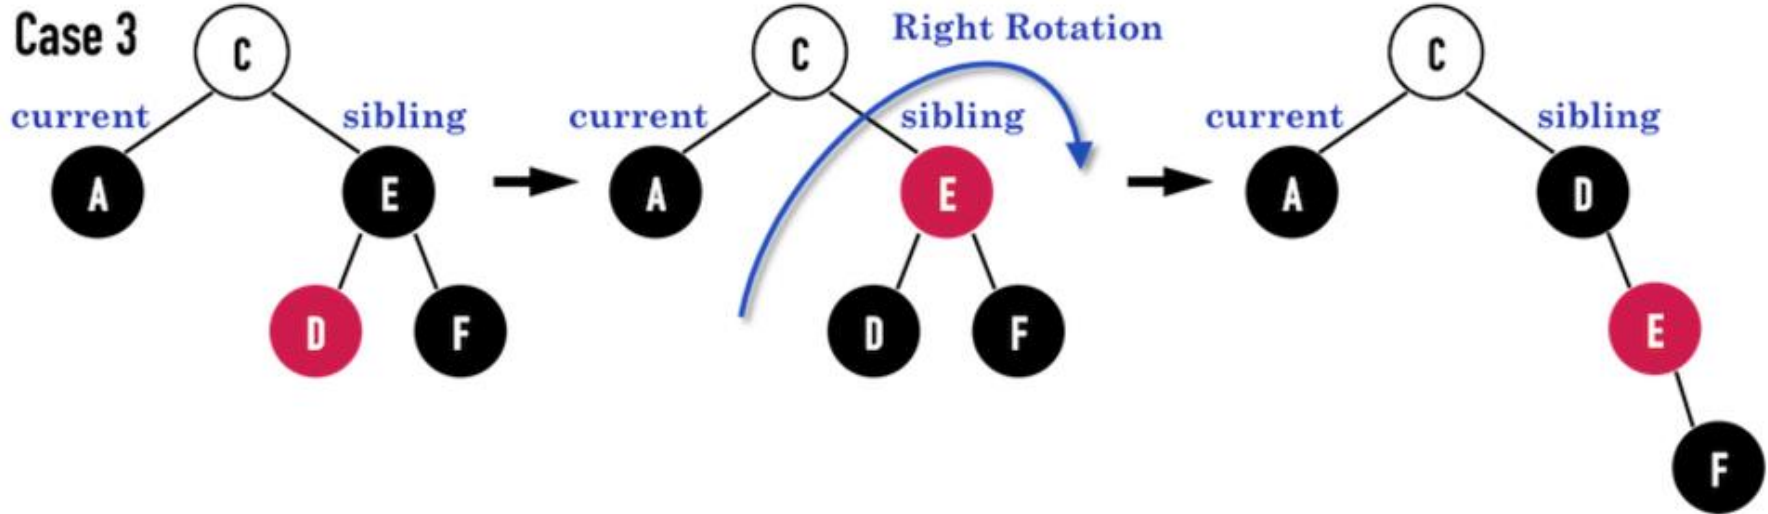
1. balance the black height of its two sub-trees,
2. maintain the black height of sub-tree C.



If current (c) is black, then we have to start over again. Note that $bh(c)$ and $bh(I)$ are different.


PS: go to case 2 in this example, but can be cases 3 and 4.

Case 3



1. sibling -> left child turns black
2. sibling turns red
3. right rotation at sibling


The goal is change case 3 to case 4.



1. exchange the color of sibling (E) and parent (C)
2. sibling right child (F) turns black
3. left rotate at parent (c)
4. current -> parent

The goal is to balance the black heights of sub-trees (c) after A is deleted.

An example



RB-DELETE-FIXUP(T, x)

■ RB-DELETE-FIXUP

```
1  while  $x \neq \text{root}[T]$  and  $\text{color}[x] = \text{BLACK}$ 
2      do if  $x = \text{left}[p[x]]$ 
3          then  $w \leftarrow \text{right}[p[x]]$ 
4              if  $\text{color}[w] = \text{RED}$ 
5                  then  $\text{color}[w] \leftarrow \text{BLACK}$            Case1
6                       $\text{color}[p[x]] = \text{RED}$            Case1
7                       $\text{LEFT-ROTATE}(T, p[x])$            Case1
8                       $w \leftarrow \text{right}[p[x]]$            Case1
```


```


9      if  $color[right[w]] = \text{BLACK}$  and
       $color[right[w]] = \text{BLACK}$ 
10     then  $color[w] \leftarrow \text{RED}$  Case2
11          $x \leftarrow p[x]$  Case2
12     else if  $color[right[w]] = \text{BLACK}$ 
13         then  $color[left[w]] \leftarrow \text{BLACK}$  Case3
14              $color[w] \leftarrow \text{RED}$  Case3
15              $\text{RIGHT-ROTATE}(T, w)$  Case3
16              $w \leftarrow right[p[x]]$  Case3
17              $color[w] \leftarrow color[p[x]]$  Case4

```

18	$color[p[x]] \leftarrow \text{BLACK}$	Case4
19	$color[right[w]] \leftarrow \text{BLACK}$	Case4
20	LEFT-ROTATE($T, p[x]$)	Case4
21	$x \leftarrow root[T]$	Case4
22	else (same as then clause with “right” and “left” exchanged)	
23	$color[x] \leftarrow \text{BLACK}$	

the case in the while loop of RB-DELETE-FIXUP





Analysis

- The RB-DELETE-FIXUP takes $O(\lg n)$ time and performs at most three rotations.
- The overall time for RB-DELETE is therefore also $O(\lg n)$