



14. Augmenting Data Structures

Yu-Shuen Wang, CS, NCTU

Augmenting Data Structures


- It will suffice to augment a textbook data structure by storing additional information in it. You can then program new operations for the data structure to support the desired application. Augmenting a data structure is not always straightforward, however, since the added information must be updated and maintained by the ordinary operations on the data structure.

14.1 Dynamic order statistics

- We shall also see the **rank** of an element—its position in the linear order of the set—can likewise be determined in $O(\lg n)$ time.

- Beside the usual red-black tree fields $key[x]$, $color[x]$, $p[x]$, $left[x]$, and $right[x]$ in a node x , we have another field $size[x]$. This field contains the number of (internal) nodes in the subtree rooted at x (including x itself), that is the size of the subtree. If we define the sentinel's size to be 0, that is, we set $size[nil[T]]$ to be 0, then we have the identity
$$size[x] = size[left[x]] + size[right[x]] + 1$$

An order-statistic tree



Retrieving an element with a given rank

```
OS-SELECT( $x, i$ )  
1   $r \leftarrow \text{size}[\text{left}[x]] + 1$   
2      if  $i = r$   
3          then return  $x$   
4      else if  $i < r$   
5          then return OS-SELECT( $\text{left}[x], i$ )  
6      else return OS-SELECT( $\text{right}[x], i - r$ )
```

Time complexity : $O(\lg n)$

Determining the rank of an element

OS-RANK(T, x)

```
1   $r \leftarrow \text{size}[\text{left}[x]] + 1$ 
2   $y \leftarrow x$ 
3  while  $y \neq \text{root}[T]$ 
4      do if  $y = \text{right}[p[y]]$ 
5          then  $r \leftarrow r + \text{size}[\text{left}[p[y]]] + 1$ 
6           $y \leftarrow p[y]$ 
7  return  $r$ 
```

The running time of OS-RANK is at worst proportional to the height of the tree: $O(\lg n)$


Maintaining subtree sizes

- Referring to the code for LEFT-ROTATE(T, x) in *section 13.2* we add the following lines:

12 $size[y] \leftarrow size[x]$

13 $size[x] \leftarrow size[left[x]] + size[right[x]] + 1$

Updating subtree sizes during rotations



14.2 How to augment a data structure

- 1.Choosing an underlying data structure,
- 2.Determining additional information to be maintained in the underlying data structure,
- 3.Verifying that the additional information can be maintained for the basic modifying operations on the underlying data structure, and
- 4.Developing new operations.

Augmenting red-black trees

- **Theorem 14.1** (Augmenting a red-black tree)

Let f be a field that augments a red-black tree T of n nodes, and suppose that the contents of f for a node x can be computed using only the information in nodes x , $left[x]$, and $right[x]$, including $f[left[x]]$ and $f[right[x]]$. Then, we can maintain the values of f in all nodes of T during insertion and deletion without asymptotically affecting the $O(\lg n)$ performance of these operations.




Proof.

- The main idea of the proof is that a change to an f field in a node x propagates only to ancestors of x in the tree.

14.3 Interval trees

- We can represent an interval $[t_1, t_2]$ as an object i , with fields $low[i] = t_1$ (the **low endpoint**) and $high[i] = t_2$ (the **high endpoint**). We say that intervals i and i' **overlap** if $i \cap i' \neq \emptyset$, that is, if $low[i] \leq high[i']$ and $low[i'] \leq high[i]$. Any two intervals i and i' satisfy the **interval trichotomy**; that exactly one of the following three properties holds:
 - i and i' overlap,
 - i is to the left of i' (i.e., $high[i] < low[i']$),
 - i is to the right of i' (i.e., $high[i'] < low[i]$)

The interval trichotomy for two closed intervals i and i'




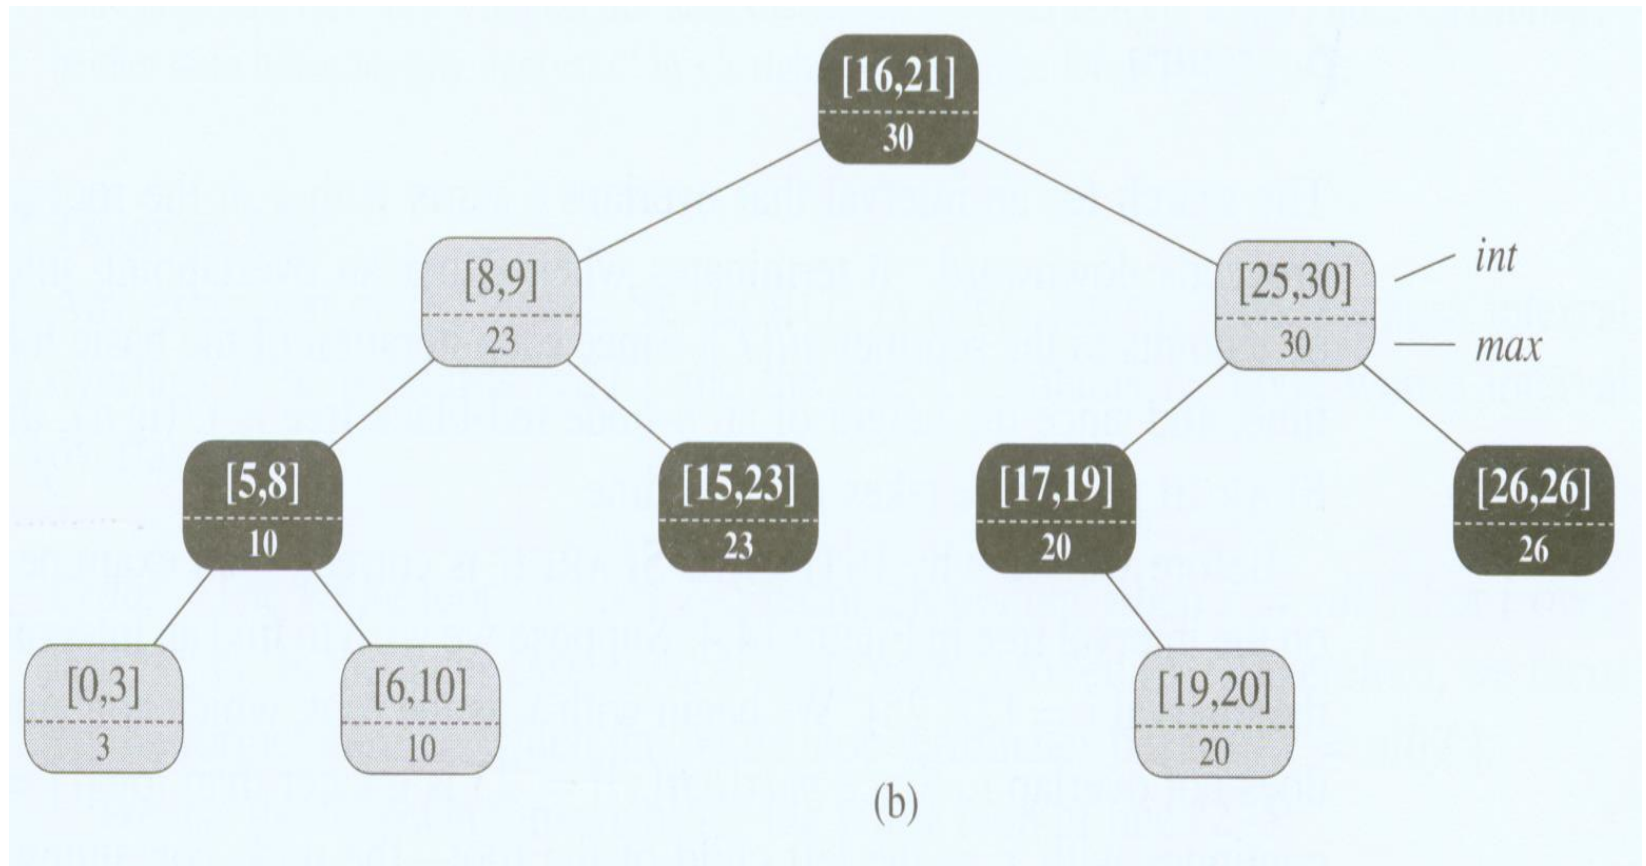
- An interval tree is a red-black tree that maintains a dynamic set of elements, with each element x containing an interval $\text{int}[x]$, and the key of x is the low endpoint, $\text{low}[\text{int}[x]]$, of the interval.

Operations

- Interval trees support the following operations.
 - INTERVAL-INSERT(T, x)
 - INTERVAL-DELETE(T, x)
 - INTERVAL-SEARCH(T, i)

An interval tree





(b)

Design of an interval tree

- Step 1: underlying data structure
 - Red-black tree
- Step 2: Additional information
 - Each node x contains a value $max[x]$, which is the maximum value of any interval endpoint stored in the subtree rooted at x .
- Step 3: Maintaining the information
 - We determine $max[x]$ given interval $int[x]$ and the max values of node x 's children:
 $max[x] = max(high[int[x]], max[left[x]], max[right[x]]).$
 - Thus, by Theorem 14.1, insertion and deletion run in $O(\lg n)$ time.

- Step 4: Develop new operations
 - The only new operation we need is INTERVAL-SEARCH(T, i), which finds a node in tree T whose interval overlaps interval i . If there is no interval that overlaps i in the tree, a pointer to the sentinel $nil[T]$ is returned.

■ INTERVAL-SEARCH(T, i)

```
1   $x \leftarrow \text{root}[T]$ 
2  while  $x \neq \text{nil}[T]$  and  $i$  does not overlap  $\text{int}[x]$ 
3      do if  $\text{left}[x] \neq \text{nil}[T]$  and  $\text{max}[\text{left}[x]] \geq \text{low}[i]$ 
4          then  $x \leftarrow \text{left}[x]$ 
5          else  $x \leftarrow \text{right}[x]$ 
6  return  $x$ 
```

Theorem 14.2

Any execution of $\text{INTERVAL-SEARCH}(T, i)$ either returns a node whose interval overlaps i , or it returns $\text{nil}[T]$ and the tree T contain node whose interval overlaps i .